

# Data Exploration

## Spark Session Initialization

```
In [ ]: from pyspark.sql import SparkSession

spark = SparkSession.builder \
    .appName("Data Exploration") \
    .config('spark.jars.packages', 'mysql:mysql-connector-java:8.0.32') \
    .getOrCreate()
```

## Data Loading

```
In [ ]: mysql_host = "sql8.freemysql.com"
mysql_port = "3306"
mysql_database = "sql8696474"
mysql_username = "sql8696474"
mysql_password = "2gVPjJi7xV"
mysql_table = "jendouba_sales"

jdbc_url = f"jdbc:mysql://{mysql_host}:{mysql_port}/{mysql_database}"

mysql_properties = {
    "user": mysql_username,
    "password": mysql_password,
    "driver": "com.mysql.cj.jdbc.Driver"
}

df = spark.read.jdbc(url=jdbc_url, table=mysql_table, properties=mysql_properties)
```

## Displaying the first few rows of the DataFrame

```
In [ ]: df.show(n=3, truncate=True)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time
15526	3	Jendouba Ain Drahem	Blueberry Muffin	12.0	3.8	45.6	livraison	online	2024-02-19 13:31:42
15526	3	Jendouba Ain Drahem	Bear Claw	20.0	6.8	136.0	livraison	online	2024-02-19 13:31:42
15526	3	Jendouba Ain Drahem	Baguette	14.0	2.0	28.0	direct	card	2024-02-19 13:31:42

only showing top 3 rows

## Displaying the schema of the DataFrame

```
In [ ]: df.printSchema()
```

```
root
 |-- client_id: integer (nullable = true)
 |-- pos_id: integer (nullable = true)
 |-- pos_name: string (nullable = true)
 |-- article: string (nullable = true)
 |-- quantity: double (nullable = true)
 |-- price: double (nullable = true)
 |-- total: double (nullable = true)
 |-- sale_type: string (nullable = true)
 |-- payment_mode: string (nullable = true)
 |-- sale_time: timestamp (nullable = true)
```

## Calculating Descriptive Statistics

```
In [ ]: df.describe("total").show()
```

summary	total
count	9672
mean	67.66650124069496
stddev	51.39601950189552
min	1.5
max	240.0

## Calculation of Correlation

```
In [ ]: from pyspark.sql.functions import corr

df.select(corr("quantity", "total")).show()
```

```
+-----+
|corr(quantity, total)|
+-----+
|    0.7167740253074978|
+-----+
```

## Counting the number of rows and distinct values

### Counting the rows

```
In [ ]: df.count()

Out[ ]: 9876
```

### Distinct values

```
In [ ]: df.select("article").distinct().count()

Out[ ]: 23

In [ ]: df.select("article").distinct().show(23)
```

```
+-----+
|      article|
+-----+
|    Cream Puff|
|         Muffin|
|    Pecan Pie|
|    Napoleon|
|         Scone|
|         null|
|    Cherry Pie|
|    Cheesecake|
| Apple Turnover|
|    Baguette|
|    Palmier|
|    Croissant|
|    Bear Claw|
|Chocolate Eclair|
|         Cupcake|
|    Key Lime Tart|
| Red Velvet Cake|
|    Lemon Bar|
|    Strudel|
|Blueberry Muffin|
|    Cinnamon Roll|
|    Fruit Tart|
|    Danish Pastry|
+-----+
```

### Distinct rows

```
In [ ]: Distinct_Df = df.distinct()

In [ ]: df.count()

Out[ ]: 9876

In [ ]: Distinct_Df.count()

Out[ ]: 9410
```

## Search and Removal of Duplicates and Missing Values

### Search and Removal of Duplicates

#### Searching for Duplicate Values

```
In [ ]: from pyspark.sql.functions import col

df.groupBy("article").count().where(col("count") > 1).show(5)

+-----+-----+
|  article|count|
+-----+-----+
|Cream Puff|  429|
|   Muffin|  427|
| Pecan Pie|  411|
|  Napoleon|  418|
|    Scone|  467|
+-----+-----+
only showing top 5 rows
```

```
In [ ]: from pyspark.sql.functions import desc, col

duplicate_values = df.groupBy("pos_id","pos_name","article","sale_time").count().where(col("count") > 1).orderBy(desc("count")).show(5)

+-----+-----+-----+-----+-----+-----+
|pos_id|      pos_name|article|      sale_time|count|
+-----+-----+-----+-----+-----+
|    4| Jendouba_Tabarka|   null|2024-04-03 22:59:54|   54|
|    3|Jendouba_Ain_Drahem|   null|2024-04-03 22:59:54|   54|
|    2| Jendouba_Bousalem|   null|2024-04-03 22:59:54|   48|
|    1| Jendouba_Center|   null|2024-04-03 22:59:54|   48|
|    3|Jendouba_Ain_Drahem|Palmier|2024-04-03 22:53:30|    7|
+-----+-----+-----+-----+-----+
only showing top 5 rows
```

Removing Duplicate Values

```
In [ ]: new_df = df.dropDuplicates(["pos_id","pos_name","article","sale_time"])

In [ ]: df.count()

Out[ ]: 9876

In [ ]: new_df.count()

Out[ ]: 8894
```

Searching for Duplicate Rows

```
In [ ]: from pyspark.sql.functions import col

df.groupBy(df.columns).count().where(col("count") > 1).show(5)

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|client_id|pos_id|      pos_name|  article|quantity|price|total|sale_type|payment_mode|      sale_time|count|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|   29402|    3|Jendouba_Ain_Drahem|Cherry Pie|    5.0|  9.5| 47.5|livraison|    online|2023-06-26 16:39:41|    2|
|    null|    2| Jendouba_Bousalem|    null|    null| null| null|livraison|    online|2024-04-03 22:59:54|   22|
|   72249|    4| Jendouba_Tabarka| Croissant|   10.0|  1.5| 15.0|livraison|    online|2023-03-10 03:01:41|    2|
|    2502|    3|Jendouba_Ain_Drahem| Croissant|   14.0|  1.5| 21.0|  direct|    cash|2023-10-25 17:55:03|    2|
|   71085|    4| Jendouba_Tabarka|Cherry Pie|   20.0|  9.5|190.0|  direct|    cash|2023-04-26 17:12:23|    2|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 5 rows
```

```
In [ ]: from pyspark.sql.functions import col

df.groupBy(df.columns).count().where(col("count") > 1).count()

Out[ ]: 285
```

Removing Duplicate Rows

```
In [ ]: new_df = df.dropDuplicates()

In [ ]: df.count()

Out[ ]: 9876

In [ ]: new_df.count()

Out[ ]: 9410
```

Searching for Missing Values

```
In [ ]: from pyspark.sql.functions import col

df.filter(col("article").isNull()).show(5)

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|client_id|pos_id|      pos_name|article|quantity|price|total|sale_type|payment_mode|      sale_time|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|    null|    3|Jendouba_Ain_Drahem|   null|    null| null| null|livraison|    online|2024-04-03 22:59:54|
|    null|    3|Jendouba_Ain_Drahem|   null|    null| null| null|livraison|    online|2024-04-03 22:59:54|
|    null|    3|Jendouba_Ain_Drahem|   null|    null| null| null|livraison|    online|2024-04-03 22:59:54|
|    null|    3|Jendouba_Ain_Drahem|   null|    null| null| null|livraison|    online|2024-04-03 22:59:54|
|    null|    2| Jendouba_Bousalem|   null|    null| null| null|livraison|    online|2024-04-03 22:59:54|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 5 rows

In [ ]: from pyspark.sql.functions import col

df.filter(col("article").isNull()).count()

Out[ ]: 204
```

## Deleting Rows Containing Missing Values

### Deleting Rows Containing Null Values

#### Method 1 dropna

```
In [ ]: cleaned_df = df.dropna(how="any")

In [ ]: cleaned_df.filter(col("article").isNull()).show()

+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|client_id|pos_id|pos_name|article|quantity|price|total|sale_type|payment_mode|sale_time|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
```

#### Method 2 na.drop

```
In [ ]: cleaned_df = df.na.drop(how="any")

In [ ]: cleaned_df.filter(col("article").isNull()).show()

+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|client_id|pos_id|pos_name|article|quantity|price|total|sale_type|payment_mode|sale_time|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
```

### Deleting Rows Containing Null Values in Specific Cloumns "Article" and "Quantity"

```
In [ ]: cleaned_df = df.na.drop(subset=["article", "quantity"])

In [ ]: cleaned_df.filter(col("article").isNull()).show()

+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|client_id|pos_id|pos_name|article|quantity|price|total|sale_type|payment_mode|sale_time|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
```

## Replacing Missing (Null) Values with a Specific Value

### Method 1: na.fill

#### Test 1 For Numerical Data

```
In [ ]: filled_df = df.na.fill(0)

In [ ]: filled_df.filter(col("article").isNull()).show(5)

+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|client_id|pos_id|pos_name|article|quantity|price|total|sale_type|payment_mode|sale_time|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|0|3|Jendouba_Ain_Drahem|null|0.0|0.0|0.0|livraison|online|2024-04-03 22:59:54|
|0|3|Jendouba_Ain_Drahem|null|0.0|0.0|0.0|livraison|online|2024-04-03 22:59:54|
|0|3|Jendouba_Ain_Drahem|null|0.0|0.0|0.0|livraison|online|2024-04-03 22:59:54|
|0|3|Jendouba_Ain_Drahem|null|0.0|0.0|0.0|livraison|online|2024-04-03 22:59:54|
|0|2|Jendouba_Bousalem|null|0.0|0.0|0.0|livraison|online|2024-04-03 22:59:54|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 5 rows
```

#### Test 2 For Numerical And Categorical Data

```
In [ ]: filled_df = df.na.fill("0")

In [ ]: filled_df.filter(col("article").isNull()).show(5)

+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|client_id|pos_id|pos_name|article|quantity|price|total|sale_type|payment_mode|sale_time|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
```

### Method 2: fillna

```
In [ ]: filled_df = df.fillna(0)
```

## Column Selection

```
In [ ]: df.select("Total").show(n=10, truncate=True)
```

```
+-----+
|Total|
+-----+
| 45.6|
|136.0|
| 28.0|
| 90.0|
| 49.0|
| 15.0|
|190.0|
| 64.0|
|100.0|
| 72.0|
+-----+
```

only showing top 10 rows

## Column Renaming

### Method 1: Using the "withColumnRenamed()" method

```
In [ ]: renamed_column_df = df.withColumnRenamed("article", "product")
renamed_column_df.show(5)
```

client_id	pos_id	pos_name	product	quantity	price	total	sale_type	payment_mode	sale_time
15526	3	Jendouba_Ain_Drahem	Blueberry Muffin	12.0	3.8	45.6	livraison	online	2024-02-19 13:31:42
15526	3	Jendouba_Ain_Drahem	Bear Claw	20.0	6.8	136.0	livraison	online	2024-02-19 13:31:42
15526	3	Jendouba_Ain_Drahem	Baguette	14.0	2.0	28.0	direct	card	2024-02-19 13:31:42
15526	3	Jendouba_Ain_Drahem	Lemon Bar	15.0	6.0	90.0	livraison	online	2024-02-19 13:31:42
75376	3	Jendouba_Ain_Drahem	Muffin	14.0	3.5	49.0	direct	card	2023-03-08 18:29:41

only showing top 5 rows

### Method 2: Using the "alias()" method

```
In [ ]: aliased_column_df = df.select("pos_id", "pos_name", df["article"].alias("product"), "quantity", "price", "total", "sale_type", "payment_mode")
aliased_column_df.show(5)
```

pos_id	pos_name	product	quantity	price	total	sale_type	payment_mode	sale_time
3	Jendouba_Ain_Drahem	Blueberry Muffin	12.0	3.8	45.6	livraison	online	2024-02-19 13:31:42
3	Jendouba_Ain_Drahem	Bear Claw	20.0	6.8	136.0	livraison	online	2024-02-19 13:31:42
3	Jendouba_Ain_Drahem	Baguette	14.0	2.0	28.0	direct	card	2024-02-19 13:31:42
3	Jendouba_Ain_Drahem	Lemon Bar	15.0	6.0	90.0	livraison	online	2024-02-19 13:31:42
3	Jendouba_Ain_Drahem	Muffin	14.0	3.5	49.0	direct	card	2023-03-08 18:29:41

only showing top 5 rows

## Creating New Columns

### Example 1: Creating a New column by multiplying two existing columns

```
In [ ]: from pyspark.sql.functions import col

new_df1 = df.withColumn("total_price", col("quantity") * col("price"))
new_df1.show(5)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time	total_price
15526	3	Jendouba_Ain_Drahem	Blueberry Muffin	12.0	3.8	45.6	livraison	online	2024-02-19 13:31:42	45.599999999999994
15526	3	Jendouba_Ain_Drahem	Bear Claw	20.0	6.8	136.0	livraison	online	2024-02-19 13:31:42	136.0
15526	3	Jendouba_Ain_Drahem	Baguette	14.0	2.0	28.0	direct	card	2024-02-19 13:31:42	28.0
15526	3	Jendouba_Ain_Drahem	Lemon Bar	15.0	6.0	90.0	livraison	online	2024-02-19 13:31:42	90.0
75376	3	Jendouba_Ain_Drahem	Muffin	14.0	3.5	49.0	direct	card	2023-03-08 18:29:41	49.0

only showing top 5 rows

### Example 2: Creating a New column by concatenating two existing columns

```
In [ ]: from pyspark.sql.functions import concat, lit

new_df2 = df.withColumn("Concatenated_Column", concat(df['quantity'], lit(" * "), df['price']))
new_df2.show(5)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time	Concatenated_Column
15526	3	Jendouba_Ain_Drahem	Blueberry Muffin	12.0	3.8	45.6	livraison	online	2024-02-19 13:31:42	12.0 * 3.8
15526	3	Jendouba_Ain_Drahem	Bear Claw	20.0	6.8	136.0	livraison	online	2024-02-19 13:31:42	20.0 * 6.8
15526	3	Jendouba_Ain_Drahem	Baguette	14.0	2.0	28.0	direct	card	2024-02-19 13:31:42	14.0 * 2.0
15526	3	Jendouba_Ain_Drahem	Lemon Bar	15.0	6.0	90.0	livraison	online	2024-02-19 13:31:42	15.0 * 6.0
75376	3	Jendouba_Ain_Drahem	Muffin	14.0	3.5	49.0	direct	card	2023-03-08 18:29:41	14.0 * 3.5

only showing top 5 rows

### Example 3: Creating a New Column Using Conditions "When"

```
In [ ]: from pyspark.sql.functions import when

new_df3 = df.withColumn("is_high_quantity", when(df['quantity'] > 10, 1).otherwise(0))
new_df3.show(5)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time	is_high_quantity
15526	3	Jendouba_Ain_Drahem	Blueberry Muffin	12.0	3.8	45.6	livraison	online	2024-02-19 13:31:42	1
15526	3	Jendouba_Ain_Drahem	Bear Claw	20.0	6.8	136.0	livraison	online	2024-02-19 13:31:42	1
15526	3	Jendouba_Ain_Drahem	Baguette	14.0	2.0	28.0	direct	card	2024-02-19 13:31:42	1
15526	3	Jendouba_Ain_Drahem	Lemon Bar	15.0	6.0	90.0	livraison	online	2024-02-19 13:31:42	1
75376	3	Jendouba_Ain_Drahem	Muffin	14.0	3.5	49.0	direct	card	2023-03-08 18:29:41	1

only showing top 5 rows

### Dropping Column

```
In [ ]: df_without_total = df.drop("total")
df_without_total.show(5)
```

client_id	pos_id	pos_name	article	quantity	price	sale_type	payment_mode	sale_time
15526	3	Jendouba_Ain_Drahem	Blueberry Muffin	12.0	3.8	livraison	online	2024-02-19 13:31:42
15526	3	Jendouba_Ain_Drahem	Bear Claw	20.0	6.8	livraison	online	2024-02-19 13:31:42
15526	3	Jendouba_Ain_Drahem	Baguette	14.0	2.0	direct	card	2024-02-19 13:31:42
15526	3	Jendouba_Ain_Drahem	Lemon Bar	15.0	6.0	livraison	online	2024-02-19 13:31:42
75376	3	Jendouba_Ain_Drahem	Muffin	14.0	3.5	direct	card	2023-03-08 18:29:41

only showing top 5 rows

### Value Replacement

```
In [ ]: from pyspark.sql.functions import regexp_replace

df_replace = df.withColumn("article", regexp_replace("article", " ", "_"))
df_replace.show(5)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time
15526	3	Jendouba_Ain_Drahem	Blueberry_Muffin	12.0	3.8	45.6	livraison	online	2024-02-19 13:31:42
15526	3	Jendouba_Ain_Drahem	Bear_Claw	20.0	6.8	136.0	livraison	online	2024-02-19 13:31:42
15526	3	Jendouba_Ain_Drahem	Baguette	14.0	2.0	28.0	direct	card	2024-02-19 13:31:42
15526	3	Jendouba_Ain_Drahem	Lemon_Bar	15.0	6.0	90.0	livraison	online	2024-02-19 13:31:42
75376	3	Jendouba_Ain_Drahem	Muffin	14.0	3.5	49.0	direct	card	2023-03-08 18:29:41

only showing top 5 rows

### Value Transformation

#### Transforming values to lowercase

```
In [ ]: from pyspark.sql.functions import lower

df_lower = df.withColumn("article_lower", lower(df['article']))
df_lower.show(5)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time	article_lower
15526	3	Jendouba_Ain_Drahem	Blueberry Muffin	12.0	3.8	45.6	livraison	online	2024-02-19 13:31:42	blueberry muffin
15526	3	Jendouba_Ain_Drahem	Bear Claw	20.0	6.8	136.0	livraison	online	2024-02-19 13:31:42	bear claw
15526	3	Jendouba_Ain_Drahem	Baguette	14.0	2.0	28.0	direct	card	2024-02-19 13:31:42	baguette
15526	3	Jendouba_Ain_Drahem	Lemon Bar	15.0	6.0	90.0	livraison	online	2024-02-19 13:31:42	lemon bar
75376	3	Jendouba_Ain_Drahem	Muffin	14.0	3.5	49.0	direct	card	2023-03-08 18:29:41	muffin

only showing top 5 rows

## Transforming Values to Uppercase

```
In [ ]: from pyspark.sql.functions import upper

df_upper = df.withColumn("pos_name_upper", upper(df['pos_name']))
df_upper.show(5)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time	pos_name_upper
15526	3	Jendouba_Ain_Drahem	Blueberry Muffin	12.0	3.8	45.6	livraison	online	2024-02-19 13:31:42	JENDOUBA_AIN_DRAHEM
15526	3	Jendouba_Ain_Drahem	Bear Claw	20.0	6.8	136.0	livraison	online	2024-02-19 13:31:42	JENDOUBA_AIN_DRAHEM
15526	3	Jendouba_Ain_Drahem	Baguette	14.0	2.0	28.0	direct	card	2024-02-19 13:31:42	JENDOUBA_AIN_DRAHEM
15526	3	Jendouba_Ain_Drahem	Lemon Bar	15.0	6.0	90.0	livraison	online	2024-02-19 13:31:42	JENDOUBA_AIN_DRAHEM
75376	3	Jendouba_Ain_Drahem	Muffin	14.0	3.5	49.0	direct	card	2023-03-08 18:29:41	JENDOUBA_AIN_DRAHEM

only showing top 5 rows

## Capitalizing the First Letter of Each Word (Example: Firstname Lastname)

```
In [ ]: from pyspark.sql.functions import initcap

df_initcap = df.withColumn("article_initcap", initcap(df['article']))
df_initcap.show(5)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time	article_initcap
15526	3	Jendouba_Ain_Drahem	Blueberry Muffin	12.0	3.8	45.6	livraison	online	2024-02-19 13:31:42	Blueberry Muffin
15526	3	Jendouba_Ain_Drahem	Bear Claw	20.0	6.8	136.0	livraison	online	2024-02-19 13:31:42	Bear Claw
15526	3	Jendouba_Ain_Drahem	Baguette	14.0	2.0	28.0	direct	card	2024-02-19 13:31:42	Baguette
15526	3	Jendouba_Ain_Drahem	Lemon Bar	15.0	6.0	90.0	livraison	online	2024-02-19 13:31:42	Lemon Bar
75376	3	Jendouba_Ain_Drahem	Muffin	14.0	3.5	49.0	direct	card	2023-03-08 18:29:41	Muffin

only showing top 5 rows

## Removing Leading and Trailing Spaces

```
In [ ]: from pyspark.sql.functions import trim

df_trimmed = df.withColumn("pos_name_trimmed", trim(df['pos_name']))
df_trimmed.show(5)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time	pos_name_trimmed
15526	3	Jendouba_Ain_Drahem	Blueberry Muffin	12.0	3.8	45.6	livraison	online	2024-02-19 13:31:42	Jendouba_Ain_Drahem
15526	3	Jendouba_Ain_Drahem	Bear Claw	20.0	6.8	136.0	livraison	online	2024-02-19 13:31:42	Jendouba_Ain_Drahem
15526	3	Jendouba_Ain_Drahem	Baguette	14.0	2.0	28.0	direct	card	2024-02-19 13:31:42	Jendouba_Ain_Drahem
15526	3	Jendouba_Ain_Drahem	Lemon Bar	15.0	6.0	90.0	livraison	online	2024-02-19 13:31:42	Jendouba_Ain_Drahem
75376	3	Jendouba_Ain_Drahem	Muffin	14.0	3.5	49.0	direct	card	2023-03-08 18:29:41	Jendouba_Ain_Drahem

only showing top 5 rows

## Extracting Substrings from a Column

```
In [ ]: from pyspark.sql.functions import substring

df_substring = df.withColumn("article_substring", substring(df['sale_time'], 1, 4))
df_substring.show(5)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time	article_substring
15526	3	Jendouba_Ain_Drahem	Blueberry Muffin	12.0	3.8	45.6	livraison	online	2024-02-19 13:31:42	2024
15526	3	Jendouba_Ain_Drahem	Bear Claw	20.0	6.8	136.0	livraison	online	2024-02-19 13:31:42	2024
15526	3	Jendouba_Ain_Drahem	Baguette	14.0	2.0	28.0	direct	card	2024-02-19 13:31:42	2024
15526	3	Jendouba_Ain_Drahem	Lemon Bar	15.0	6.0	90.0	livraison	online	2024-02-19 13:31:42	2024
75376	3	Jendouba_Ain_Drahem	Muffin	14.0	3.5	49.0	direct	card	2023-03-08 18:29:41	2023

only showing top 5 rows

## Extracting Date and Time Components

### Adding, Subtracting Days

```
In [ ]: from pyspark.sql.functions import date_add, date_sub

date_df = df.withColumn('DateWithDay', date_add('sale_time', 1))
date_df.show(3)
```



client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time	DateWithDay
15526	3	Jendouba_Ain_Drahem	Blueberry Muffin	12.0	3.8	45.6	livraison	online	2024-02-19 13:31:42	2024-02-20
15526	3	Jendouba_Ain_Drahem	Bear Claw	20.0	6.8	136.0	livraison	online	2024-02-19 13:31:42	2024-02-20
15526	3	Jendouba_Ain_Drahem	Baguette	14.0	2.0	28.0	direct	card	2024-02-19 13:31:42	2024-02-20

only showing top 3 rows

### Extracting the Month, Year, Quarter

```
In [ ]: from pyspark.sql.functions import month, year, quarter

month_df = df.withColumn('Month', month('sale_time'))
month_df.show(3)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time	Month
15526	3	Jendouba_Ain_Drahem	Blueberry Muffin	12.0	3.8	45.6	livraison	online	2024-02-19 13:31:42	2
15526	3	Jendouba_Ain_Drahem	Bear Claw	20.0	6.8	136.0	livraison	online	2024-02-19 13:31:42	2
15526	3	Jendouba_Ain_Drahem	Baguette	14.0	2.0	28.0	direct	card	2024-02-19 13:31:42	2

only showing top 3 rows

### Extracting the Day of the Month

```
In [ ]: from pyspark.sql.functions import dayofmonth

day_df = df.withColumn('Day_Number', dayofmonth('sale_time'))
day_df.show(2)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time	Day_Number
15526	3	Jendouba_Ain_Drahem	Blueberry Muffin	12.0	3.8	45.6	livraison	online	2024-02-19 13:31:42	19
15526	3	Jendouba_Ain_Drahem	Bear Claw	20.0	6.8	136.0	livraison	online	2024-02-19 13:31:42	19

only showing top 2 rows

### Extracting the Day of the Week Name

```
In [ ]: from pyspark.sql.functions import date_format

day_name_df = df.withColumn('Day_Name', date_format('sale_time', 'EEEE'))
day_name_df.show(3)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time	Day_Name
15526	3	Jendouba_Ain_Drahem	Blueberry Muffin	12.0	3.8	45.6	livraison	online	2024-02-19 13:31:42	Monday
15526	3	Jendouba_Ain_Drahem	Bear Claw	20.0	6.8	136.0	livraison	online	2024-02-19 13:31:42	Monday
15526	3	Jendouba_Ain_Drahem	Baguette	14.0	2.0	28.0	direct	card	2024-02-19 13:31:42	Monday

only showing top 3 rows

### Extracting the Hour, Minute, second

```
In [ ]: from pyspark.sql.functions import hour, minute, second

hour_df = df.withColumn('Hour', hour('sale_time'))
hour_df.show(3)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time	Hour
15526	3	Jendouba_Ain_Drahem	Blueberry Muffin	12.0	3.8	45.6	livraison	online	2024-02-19 13:31:42	13
15526	3	Jendouba_Ain_Drahem	Bear Claw	20.0	6.8	136.0	livraison	online	2024-02-19 13:31:42	13
15526	3	Jendouba_Ain_Drahem	Baguette	14.0	2.0	28.0	direct	card	2024-02-19 13:31:42	13

only showing top 3 rows

## Data Filtering and Sorting

### Data Filtering

```
In [ ]: from pyspark.sql.functions import col
```

#### Data Filtering Using the filter() Function

```
In [ ]: filter_1 = df.filter(col('quantity') > 5)
filter_1.show(5)
```



client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time
15526	3	Jendouba_Ain_Drahem	Blueberry Muffin	12.0	3.8	45.6	livraison	online	2024-02-19 13:31:42
15526	3	Jendouba_Ain_Drahem	Bear Claw	20.0	6.8	136.0	livraison	online	2024-02-19 13:31:42
15526	3	Jendouba_Ain_Drahem	Baguette	14.0	2.0	28.0	direct	card	2024-02-19 13:31:42
15526	3	Jendouba_Ain_Drahem	Lemon Bar	15.0	6.0	90.0	livraison	online	2024-02-19 13:31:42
75376	3	Jendouba_Ain_Drahem	Muffin	14.0	3.5	49.0	direct	card	2023-03-08 18:29:41

only showing top 5 rows

Data Filtering Using the where() Function

```
In [ ]: filter_2 = df.where(col('payment_mode') == "online")
filter_2.show(5)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time
15526	3	Jendouba_Ain_Drahem	Blueberry Muffin	12.0	3.8	45.6	livraison	online	2024-02-19 13:31:42
15526	3	Jendouba_Ain_Drahem	Bear Claw	20.0	6.8	136.0	livraison	online	2024-02-19 13:31:42
15526	3	Jendouba_Ain_Drahem	Lemon Bar	15.0	6.0	90.0	livraison	online	2024-02-19 13:31:42
75376	3	Jendouba_Ain_Drahem	Apple Turnover	3.0	5.0	15.0	livraison	online	2023-03-08 18:29:41
2365	2	Jendouba_Bousalem	Strudel	8.0	8.0	64.0	livraison	online	2023-02-19 12:06:11

only showing top 5 rows

Filtering with the Logical Operator "&" (AND)

```
In [ ]: filter_3 = df.filter((col('sale_type') == 'direct') & (col('payment_mode') == 'card'))
filter_3.show(5)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time
15526	3	Jendouba_Ain_Drahem	Baguette	14.0	2.0	28.0	direct	card	2024-02-19 13:31:42
75376	3	Jendouba_Ain_Drahem	Muffin	14.0	3.5	49.0	direct	card	2023-03-08 18:29:41
2365	2	Jendouba_Bousalem	Apple Turnover	20.0	5.0	100.0	direct	card	2023-02-19 12:06:11
91492	4	Jendouba_Tabarka	Cheesecake	12.0	10.5	126.0	direct	card	2024-01-13 20:01:18
91492	4	Jendouba_Tabarka	Baguette	4.0	2.0	8.0	direct	card	2024-01-13 20:01:18

only showing top 5 rows

Filtering with the Logical Operator "|" (OR)

```
In [ ]: filter_4 = df.filter((col('sale_type') == 'direct') | (col('payment_mode') == 'card'))
filter_4.show(5)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time
15526	3	Jendouba_Ain_Drahem	Baguette	14.0	2.0	28.0	direct	card	2024-02-19 13:31:42
75376	3	Jendouba_Ain_Drahem	Muffin	14.0	3.5	49.0	direct	card	2023-03-08 18:29:41
39564	2	Jendouba_Bousalem	Cherry Pie	20.0	9.5	190.0	direct	cash	2023-11-17 02:10:48
2365	2	Jendouba_Bousalem	Apple Turnover	20.0	5.0	100.0	direct	card	2023-02-19 12:06:11
2365	2	Jendouba_Bousalem	Red Velvet Cake	6.0	12.0	72.0	direct	cash	2023-02-19 12:06:11

only showing top 5 rows

Filtering with the isin() Method

```
In [ ]: Article_List = ['Croissant', 'Chocolate Eclair', 'Fruit Tart', 'Cinnamon Roll', 'Danish Pastry']

filter_5 = df.filter(col('article').isin(Article_List))
filter_5.show(5)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time
91492	4	Jendouba_Tabarka	Danish Pastry	11.0	6.5	71.5	direct	card	2024-01-13 20:01:18
19413	1	Jendouba_Center	Croissant	13.0	1.5	19.5	livraison	online	2024-02-13 19:35:23
13685	1	Jendouba_Center	Danish Pastry	7.0	6.5	45.5	livraison	online	2023-03-20 17:35:06
13685	1	Jendouba_Center	Cinnamon Roll	14.0	4.0	56.0	livraison	online	2023-03-20 17:35:06
96107	2	Jendouba_Bousalem	Cinnamon Roll	14.0	4.0	56.0	direct	card	2024-02-29 23:34:19

only showing top 5 rows

Filtering with the Logical Operator "~" (NOT)

```
In [ ]: Article_List = ['Croissant', 'Chocolate Eclair', 'Fruit Tart', 'Cinnamon Roll', 'Danish Pastry']

filter_6 = df.filter(~col('article').isin(Article_List))
filter_6.show(5)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time
15526	3	Jendouba_Ain_Drahem	Blueberry Muffin	12.0	3.8	45.6	livraison	online	2024-02-19 13:31:42
15526	3	Jendouba_Ain_Drahem	Bear Claw	20.0	6.8	136.0	livraison	online	2024-02-19 13:31:42
15526	3	Jendouba_Ain_Drahem	Baguette	14.0	2.0	28.0	direct	card	2024-02-19 13:31:42
15526	3	Jendouba_Ain_Drahem	Lemon Bar	15.0	6.0	90.0	livraison	online	2024-02-19 13:31:42
75376	3	Jendouba_Ain_Drahem	Muffin	14.0	3.5	49.0	direct	card	2023-03-08 18:29:41

only showing top 5 rows

Filtering with the like() Method

```
In [ ]: filter_7 = df.filter(col('article').like("Apple %"))
filter_7.show(5)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time
75376	3	Jendouba_Ain_Drahem	Apple Turnover	3.0	5.0	15.0	livraison	online	2023-03-08 18:29:41
2365	2	Jendouba_Bousalem	Apple Turnover	20.0	5.0	100.0	direct	card	2023-02-19 12:06:11
96063	4	Jendouba_Tabarka	Apple Turnover	7.0	5.0	35.0	livraison	online	2023-10-20 07:35:48
75660	4	Jendouba_Tabarka	Apple Turnover	20.0	5.0	100.0	livraison	online	2023-05-18 01:27:22
87457	2	Jendouba_Bousalem	Apple Turnover	7.0	5.0	35.0	livraison	online	2023-10-29 11:36:11

only showing top 5 rows

Filtering with the endswith() Method

```
In [ ]: filter_8 = df.filter(col('article').endswith("Pie"))
filter_8.show(5)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time
39564	2	Jendouba_Bousalem	Cherry Pie	20.0	9.5	190.0	direct	cash	2023-11-17 02:10:48
19413	1	Jendouba_Center	Pecan Pie	6.0	11.0	66.0	direct	cash	2024-02-13 19:35:23
10105	2	Jendouba_Bousalem	Cherry Pie	17.0	9.5	161.5	direct	card	2023-05-09 22:35:40
96107	2	Jendouba_Bousalem	Pecan Pie	5.0	11.0	55.0	direct	card	2024-02-29 23:34:19
87457	2	Jendouba_Bousalem	Pecan Pie	12.0	11.0	132.0	direct	cash	2023-10-29 11:36:11

only showing top 5 rows

Filtering with the between() Method

```
In [ ]: filter_9 = df.filter(col("quantity").between(10.0, 20.0))
filter_9.show(5)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time
15526	3	Jendouba_Ain_Drahem	Blueberry Muffin	12.0	3.8	45.6	livraison	online	2024-02-19 13:31:42
15526	3	Jendouba_Ain_Drahem	Bear Claw	20.0	6.8	136.0	livraison	online	2024-02-19 13:31:42
15526	3	Jendouba_Ain_Drahem	Baguette	14.0	2.0	28.0	direct	card	2024-02-19 13:31:42
15526	3	Jendouba_Ain_Drahem	Lemon Bar	15.0	6.0	90.0	livraison	online	2024-02-19 13:31:42
75376	3	Jendouba_Ain_Drahem	Muffin	14.0	3.5	49.0	direct	card	2023-03-08 18:29:41

only showing top 5 rows

Filtering with the isNotNull() Method

```
In [ ]: filter_10 = df.filter(col("article").isNotNull())
filter_10.show(5)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time
15526	3	Jendouba_Ain_Drahem	Blueberry Muffin	12.0	3.8	45.6	livraison	online	2024-02-19 13:31:42
15526	3	Jendouba_Ain_Drahem	Bear Claw	20.0	6.8	136.0	livraison	online	2024-02-19 13:31:42
15526	3	Jendouba_Ain_Drahem	Baguette	14.0	2.0	28.0	direct	card	2024-02-19 13:31:42
15526	3	Jendouba_Ain_Drahem	Lemon Bar	15.0	6.0	90.0	livraison	online	2024-02-19 13:31:42
75376	3	Jendouba_Ain_Drahem	Muffin	14.0	3.5	49.0	direct	card	2023-03-08 18:29:41

only showing top 5 rows

Filtering with the isNull() Method

```
In [ ]: filter_11 = df.filter(col("article").isNull())
filter_11.show(5)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time
null	3	Jendouba_Ain_Drahem	null	null	null	null	livraison	online	2024-04-03 22:59:54
null	3	Jendouba_Ain_Drahem	null	null	null	null	livraison	online	2024-04-03 22:59:54
null	3	Jendouba_Ain_Drahem	null	null	null	null	livraison	online	2024-04-03 22:59:54
null	3	Jendouba_Ain_Drahem	null	null	null	null	livraison	online	2024-04-03 22:59:54
null	2	Jendouba_Bousalem	null	null	null	null	livraison	online	2024-04-03 22:59:54

only showing top 5 rows

# Sorting Data (Ascending and Descending) with the "sort" or "orderBy" Functions

## Sorting Data in Ascending Order

### Method 1: Using the sort() Function

```
In [ ]: sorted_df = df.sort('total')
sorted_df.show(5)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time
null	3	Jendouba_Ain_Drahem	null	null	null	null	livraison	online	2024-04-03 22:59:54
null	2	Jendouba_Bousalem	null	null	null	null	direct	cash	2024-04-03 22:59:54
null	3	Jendouba_Ain_Drahem	null	null	null	null	livraison	online	2024-04-03 22:59:54
null	3	Jendouba_Ain_Drahem	null	null	null	null	livraison	online	2024-04-03 22:59:54
null	3	Jendouba_Ain_Drahem	null	null	null	null	livraison	online	2024-04-03 22:59:54

only showing top 5 rows

### Method 2: Using the orderBy() Function

```
In [ ]: sorted_df = df.orderBy('total')
sorted_df.show(5)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time
null	3	Jendouba_Ain_Drahem	null	null	null	null	livraison	online	2024-04-03 22:59:54
null	2	Jendouba_Bousalem	null	null	null	null	direct	cash	2024-04-03 22:59:54
null	3	Jendouba_Ain_Drahem	null	null	null	null	livraison	online	2024-04-03 22:59:54
null	3	Jendouba_Ain_Drahem	null	null	null	null	livraison	online	2024-04-03 22:59:54
null	3	Jendouba_Ain_Drahem	null	null	null	null	livraison	online	2024-04-03 22:59:54

only showing top 5 rows

```
In [ ]: sorted_df = df.orderBy('total', ascending=True)
sorted_df.show(5)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time
null	3	Jendouba_Ain_Drahem	null	null	null	null	livraison	online	2024-04-03 22:59:54
null	2	Jendouba_Bousalem	null	null	null	null	direct	cash	2024-04-03 22:59:54
null	3	Jendouba_Ain_Drahem	null	null	null	null	livraison	online	2024-04-03 22:59:54
null	3	Jendouba_Ain_Drahem	null	null	null	null	livraison	online	2024-04-03 22:59:54
null	3	Jendouba_Ain_Drahem	null	null	null	null	livraison	online	2024-04-03 22:59:54

only showing top 5 rows

```
In [ ]: from pyspark.sql.functions import asc

sorted_df = df.orderBy(asc('total'))
sorted_df.show(5)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time
null	3	Jendouba_Ain_Drahem	null	null	null	null	livraison	online	2024-04-03 22:59:54
null	2	Jendouba_Bousalem	null	null	null	null	direct	cash	2024-04-03 22:59:54
null	3	Jendouba_Ain_Drahem	null	null	null	null	livraison	online	2024-04-03 22:59:54
null	3	Jendouba_Ain_Drahem	null	null	null	null	livraison	online	2024-04-03 22:59:54
null	3	Jendouba_Ain_Drahem	null	null	null	null	livraison	online	2024-04-03 22:59:54

only showing top 5 rows

## Sorting Data in Descending Order

```
In [ ]: sorted_df = df.orderBy('total', ascending=False)
sorted_df.show(5)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time
46483	1	Jendouba_Center	Red Velvet Cake	20.0	12.0	240.0	livraison	online	2023-07-30 15:11:18
26491	4	Jendouba_Tabarka	Red Velvet Cake	20.0	12.0	240.0	direct	card	2023-08-11 08:49:50
62551	3	Jendouba_Ain_Drahem	Red Velvet Cake	20.0	12.0	240.0	livraison	online	2023-10-22 16:14:39
55939	3	Jendouba_Ain_Drahem	Red Velvet Cake	20.0	12.0	240.0	direct	card	2023-06-19 09:16:36
492	1	Jendouba_Center	Red Velvet Cake	20.0	12.0	240.0	direct	card	2024-01-31 17:53:07

only showing top 5 rows

```
In [ ]: from pyspark.sql.functions import desc

sorted_df = df.orderBy(desc('total'))
sorted_df.show(5)
```


client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time
46483	1	Jendouba_Center	Red Velvet Cake	20.0	12.0	240.0	livraison	online	2023-07-30 15:11:18
26491	4	Jendouba_Tabarka	Red Velvet Cake	20.0	12.0	240.0	direct	card	2023-08-11 08:49:50
62551	3	Jendouba_Ain_Drahem	Red Velvet Cake	20.0	12.0	240.0	livraison	online	2023-10-22 16:14:39
55939	3	Jendouba_Ain_Drahem	Red Velvet Cake	20.0	12.0	240.0	direct	card	2023-06-19 09:16:36
492	1	Jendouba_Center	Red Velvet Cake	20.0	12.0	240.0	direct	card	2024-01-31 17:53:07

only showing top 5 rows

## Merge two or more DataFrames

### Difference between union intersect and except :

- UNION** operator returns all the unique rows from both the left and the right query.
- UNION ALL** included the duplicates as well.
- INTERSECT** operator retrieves the common unique rows from both the left and the right query.
- EXCEPT** operator returns unique rows from the left query that aren't in the right query's results.



Let us understand these differences with examples. We will use the following 2 tables for the examples :

Id	Name	Gender
1	Mark	Male
2	Mary	Female
3	Steve	Male
3	Steve	Male

Id	Name	Gender
2	Mary	Female
3	Steve	Male
4	John	Male

UNION :

Id	Name	Gender
1	Mark	Male
2	Mary	Female
3	Steve	Male
4	John	Male

UNION ALL :

Id	Name	Gender
1	Mark	Male
2	Mary	Female
3	Steve	Male
3	Steve	Male
2	Mary	Female
3	Steve	Male
4	John	Male

INTERSECT :

Id	Name	Gender
2	Mary	Female
3	Steve	Male

EXCEPT :

Id	Name	Gender
1	Mark	Male

```
In [ ]: mysql_host = "sql8.freesqldatabase.com"
mysql_port = "3306"
mysql_database = "sql8696474"
mysql_username = "sql8696474"
mysql_password = "2gVPjJi7xV"
mysql_table = "merge_sales"

jdbc_url = f"jdbc:mysql://{mysql_host}:{mysql_port}/{mysql_database}"

mysql_properties = {
    "user": mysql_username,
    "password": mysql_password,
    "driver": "com.mysql.cj.jdbc.Driver"
}

merge_df = spark.read.jdbc(url=jdbc_url, table=mysql_table, properties=mysql_properties)
```

```
In [ ]: df.count()
```

Out[ ]: 9876

```
In [ ]: merge_df.count()
```

Out[ ]: 528

```
In [ ]: merge_df.distinct().count()
```

Out[ ]: 523

```
In [ ]: merge_df.select("pos_id", "pos_name").distinct().show()
```

pos_id	pos_name
3	Jendouba_Ain_Drahem
2	Jendouba_Bousalem
4	Jendouba_Tabarka
1	Beja_Center
10	Tunis
3	Beja_Nefza
8	Kef
4	Beja_Testour
1	Jendouba_Center
9	Bizert
2	Beja_Amdoun

## UNION of Two DataFrames

```
In [ ]: merged_df_1 = df.union(merge_df)

merged_df_1.count()
```

Out[ ]: 10404

```
In [ ]: 9876 + 528
```

Out[ ]: 10404

## UNION ALL of Two DataFrames [link](#)

```
In [ ]: merged_df_2 = df.unionAll(merge_df)

merged_df_2.count()
```

Out[ ]: 10404

## INTERSECT of Two DataFrames

```
In [ ]: merged_df_3 = df.intersect(merge_df)

merged_df_3.count()
```

Out[ ]: 72

## EXCEPT of Two DataFrames

```
In [ ]: merged_df_4 = df.exceptAll(merge_df)

merged_df_4.count()
```


Out[ ]: 9804

```
In [ ]: 9876 - 72
```

Out[ ]: 9804

# Joining Two DataFrames

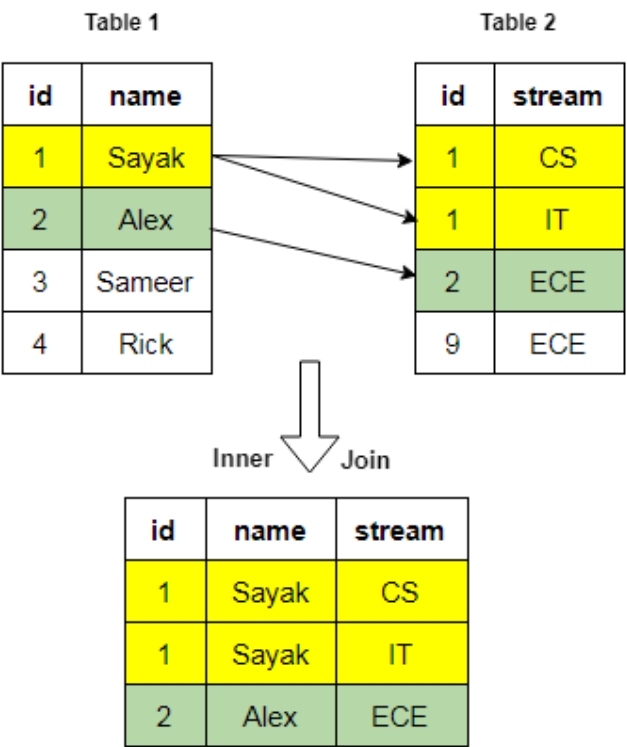
Difference between INNER, LEFT, RIGHT and FULL JOIN :



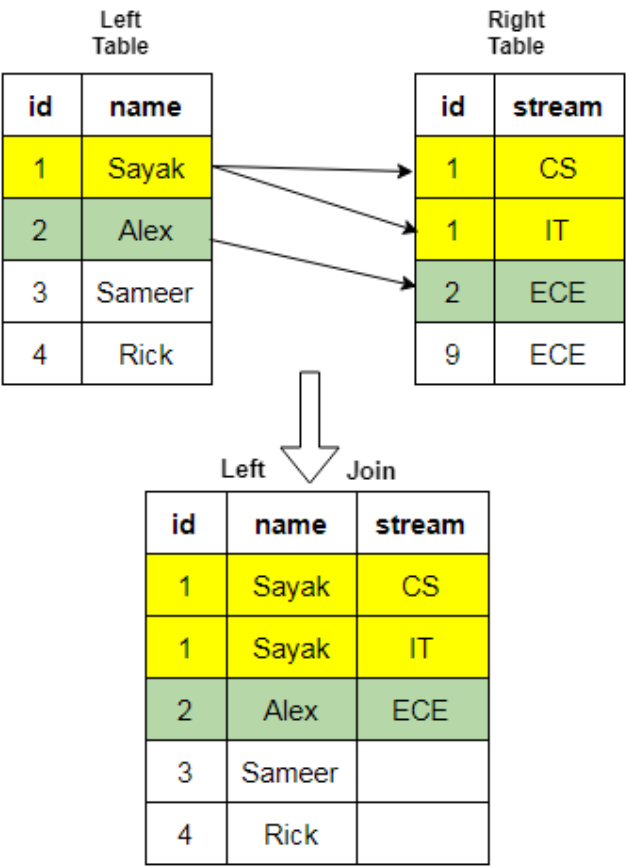
- **INNER JOIN:** Returns only the rows with matching values in both the left and right tables based on the specified condition.
- **LEFT JOIN:** Returns all rows from the left table, and the matched rows from the right table. If no match is found, NULL values are returned for the right table columns.
- **RIGHT JOIN:** Returns all rows from the right table, and the matched rows from the left table. If no match is found, NULL values are returned for the left table columns.
- **FULL JOIN:** Returns all rows from both the left and right tables, combining the results where possible. If no match is found in either table, NULL values are returned for the columns of the table with no match.

Let us understand these differences with examples :

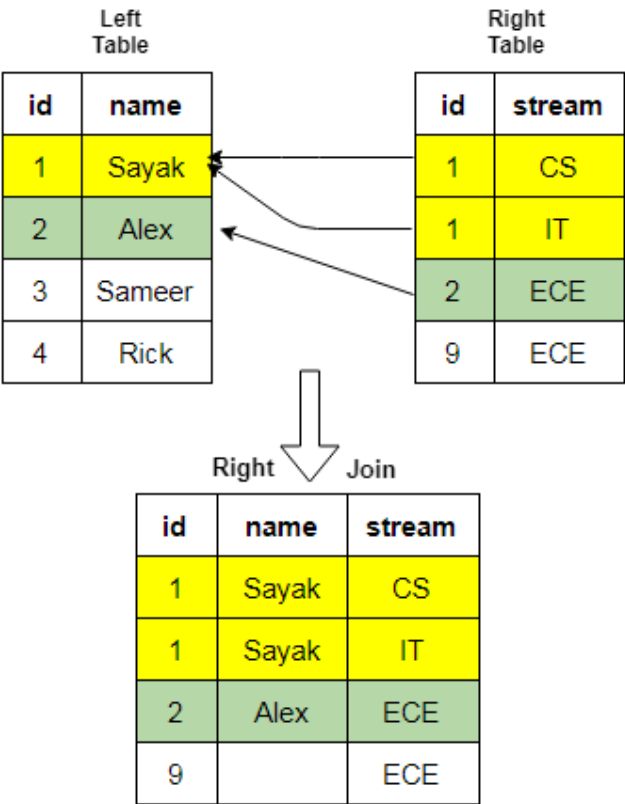
**INNER JOIN :**



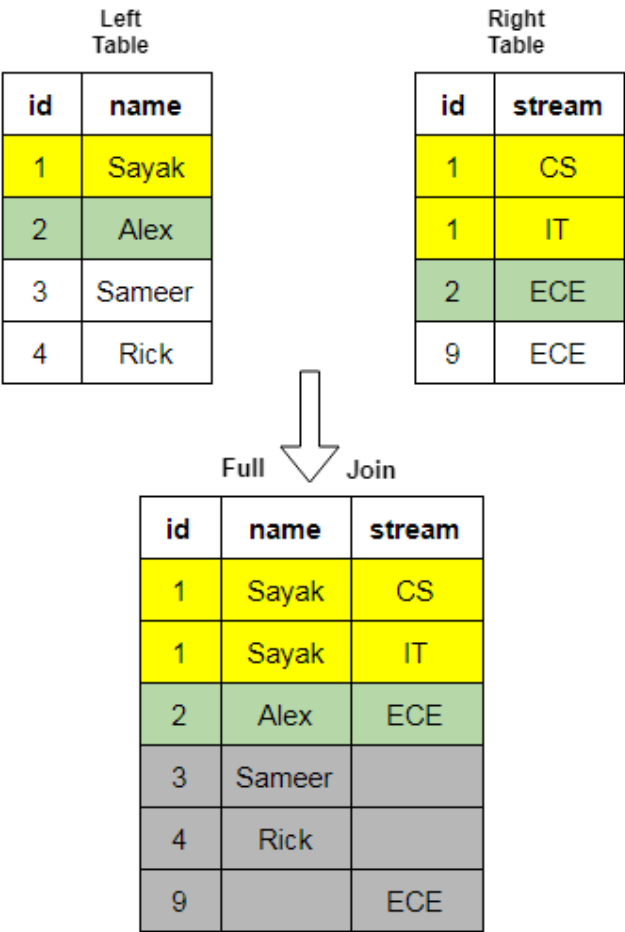
**LEFT JOIN :**



RIGHT JOIN :



FULL JOIN :



```
In [ ]: mysql_host = "sql8.freemsqldatabase.com"
mysql_port = "3306"
mysql_database = "sql8696474"
mysql_username = "sql8696474"
mysql_password = "2gVPjI7xV"
mysql_table = "join_sales"

jdbc_url = f"jdbc:mysql://{mysql_host}:{mysql_port}/{mysql_database}"

mysql_properties = { "user": mysql_username, "password": mysql_password, "driver": "com.mysql.cj.jdbc.Driver"}

join_df = spark.read.jdbc(url=jdbc_url, table=mysql_table, properties=mysql_properties)

join_df = join_df.filter(col("article") != "Muffin")
```

```
In [ ]: join_df.show(10, False)
```

article	description
Croissant	A buttery, flaky pastry originating from France.
Chocolate Eclair	A delicious pastry filled with chocolate cream and topped with chocolate icing.
Fruit Tart	A tart filled with assorted fresh fruits on top of a custard or cream filling.
Cinnamon Roll	A sweet roll served commonly in Northern Europe and North America.
Danish Pastry	A multilayered, laminated sweet pastry in the viennoiserie tradition.
Palmier	A pastry in the shape of a palm leaf or butterfly wings, made from puff pastry and sugar.
Cream Puff	A filled French pastry ball with a typically sweet and moist filling.
Apple Turnover	A pastry made by placing apple filling on a piece of dough, then folding the dough over.
Bear Claw	A sweet, yeast-raised pastry, often shaped like a bear's paw and topped with almonds.
Napoleon	A pastry made of layers of puff pastry alternating with a sweet filling, usually pastry cream.

only showing top 10 rows



clean the DataFrame df for missing values

```
In [ ]: new_df = df.dropna(how="any")
```

```
In [ ]: new_df.count()
```

Out[ ]: 9506

INNER JOIN

Joining on a Common Column

Method 1:

```
In [ ]: merged_inner = new_df.join(join_df, on='article', how='inner')

merged_inner.show(5)
```

article	client_id	pos_id	pos_name	quantity	price	total	sale_type	payment_mode	sale_time	description
Apple Turnover	75376	3	Jendouba_Ain_Drahem	3.0	5.0	15.0	livraison	online	2023-03-08 18:29:41	A pastry made by ...
Apple Turnover	2365	2	Jendouba_Bousalem	20.0	5.0	100.0	direct	card	2023-02-19 12:06:11	A pastry made by ...
Apple Turnover	96063	4	Jendouba_Tabarka	7.0	5.0	35.0	livraison	online	2023-10-20 07:35:48	A pastry made by ...
Apple Turnover	75660	4	Jendouba_Tabarka	20.0	5.0	100.0	livraison	online	2023-05-18 01:27:22	A pastry made by ...
Apple Turnover	87457	2	Jendouba_Bousalem	7.0	5.0	35.0	livraison	online	2023-10-29 11:36:11	A pastry made by ...

only showing top 5 rows

Method 2:

```
In [ ]: merged_inner = new_df.join(join_df, df.article == join_df.article, how='inner')

merged_inner.count()
```

Out[ ]: 9088

```
In [ ]: 9506 - 9088
```

Out[ ]: 418

Joining on Two Common Columns

```
In [ ]: join_condition = (df.article == join_df.article) & (df.pos_name == join_df.pos_name)

merged_inner = df.join(df_descriptions, join_condition, how='inner')
```

LEFT JOIN

```
In [ ]: merged_left = new_df.join(join_df, new_df.article == join_df.article, how='left')

merged_left.count()
```

Out[ ]: 9506

RIGHT JOIN

```
In [ ]: merged_right = new_df.join(join_df, new_df.article == join_df.article, how='right')

merged_right.count()
```

Out[ ]: 9088

```
In [ ]: 9506 - 9088
```

Out[ ]: 418

FULL JOIN

```
In [ ]: merged_full = new_df.join(join_df, new_df.article == join_df.article, how='full')

merged_full.count()
```

Out[ ]: 9506

# Aggregation, Grouping, and Window Functions

## Aggregation Functions

- **AVG()**: Calculates the average of the set of values.
- **COUNT()**: Returns the count of rows.
- **SUM()**: Calculates the arithmetic sum of the set of numeric values.
- **MAX()**: From a group of values, returns the maximum value.
- **MIN()**: From a group of values, returns the minimum value.

```
In [ ]: from pyspark.sql.functions import sum, avg,max, min, count

new_df.agg(sum('total').alias('Total_Sum'),
           avg('total').alias('Total_Mean'),
           max('total').alias('Total_Max'),
           min('total').alias('Total_Min'),
           count('total').alias('Total_Count')).show()
```

Total_Sum	Total_Mean	Total_Max	Total_Min	Total_Count
642509.0000000019	67.58983799705469	240.0	1.5	9506

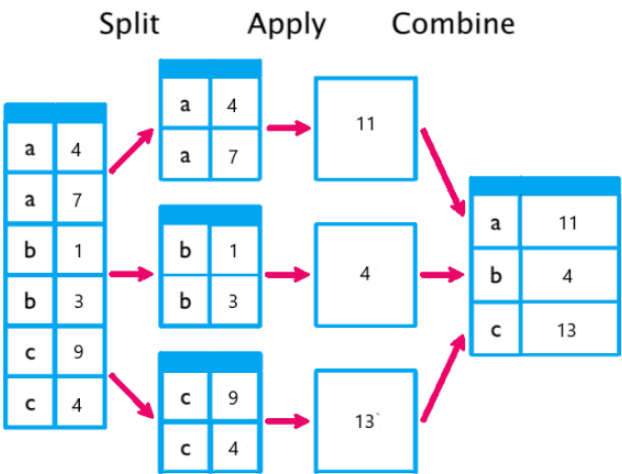
## Grouping Functions

The Group By statement is used to group together any rows of a column with the same value stored in them, based on a function specified in the statement. Generally, these functions are one of the aggregate functions such as MAX() and SUM().

The Group By statement uses the split-apply-combine strategy:

- **Split**: The different groups are split with their values.
- **Apply**: The aggregate function is applied to the values of these groups.
- **Combine**: The values are combined in a single row.

The Group By function is typically used when you want to apply multiple aggregation operations on different columns at the same time, or when you want to rename the aggregated columns.



### Example 1:

```
In [ ]: new_df.groupBy('article').sum('total').show(5)
```

article	sum(total)
Cream Puff	40077.0
Muffin	15540.0
Pecan Pie	46948.0
Napoleon	33527.599999999984
Scone	13378.399999999992

only showing top 5 rows

### Example 2:

```
In [ ]: new_df.groupBy('article') \
        .agg(
            sum('total').alias('Total_Sum'),
            count('total').alias('Total_Count')) \
        .show(5)
```

article	Total_Sum	Total_Count
Cream Puff	40077.0	422
Muffin	15540.0	418
Pecan Pie	46948.0	405
Napoleon	33527.599999999984	407
Scone	13378.399999999992	456

only showing top 5 rows

Example 3:

```
In [ ]: from pyspark.sql.functions import sum, count

new_df.groupBy('article','pos_name') \
    .agg(
        sum('total').alias('Total_Sum'),
        count('total').alias('Total_Count')) \
    .show(5)
```

article	pos_name	Total_Sum	Total_Count
Palmier	Jendouba_Bousalem	3612.0	108
Key Lime Tart	Jendouba_Center	10791.199999999999	123
Cupcake	Jendouba_Tabarka	5301.0	105
Key Lime Tart	Jendouba_Bousalem	7609.599999999996	102
Chocolate Eclair	Jendouba_Bousalem	8430.0	109

only showing top 5 rows

Pivot Tables

The column whose distinct values become new columns.

```
In [ ]: new_df.groupBy('article') \
    .pivot('sale_type') \
    .agg(sum('total').alias('Total_Sum')) \
    .show(5)
```

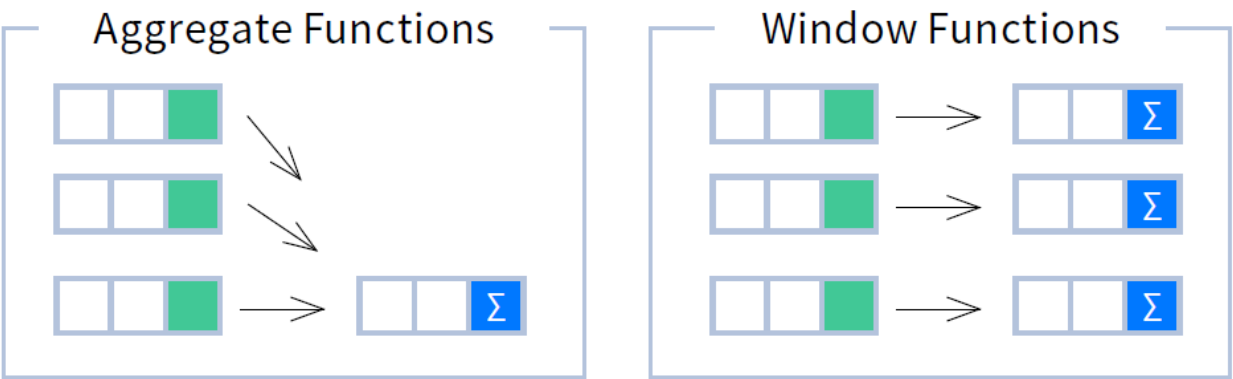
article	direct	livraison
Cream Puff	20385.0	19692.0
Muffin	8368.5	7171.5
Napoleon	18312.200000000004	15215.399999999992
Pecan Pie	21538.0	25410.0
Scone	6775.999999999996	6602.4

only showing top 5 rows

Window Functions

Aggregate Functions vs Window Functions

Unlike aggregate functions, window functions do not collapse rows.



```
In [ ]: from pyspark.sql.window import Window
from pyspark.sql.functions import sum, count
```

GroupBy Example

```
In [ ]: new_df.groupBy('article') \
    .agg(
        sum('total').alias('Total_Sum'),
        count('total').alias('Total_Count')) \
    .orderBy("Total_Count") \
    .show(5)
```

article	Total_Sum	Total_Count
Chocolate Eclair	29970.0	399
Pecan Pie	46948.0	405
Napoleon	33527.599999999998	407
Cupcake	20052.0	412
Muffin	15540.0	418

only showing top 5 rows

Window.partitionBy Example

```
In [ ]: window = Window.partitionBy('article')

windowed_df = new_df.withColumn('Total_Sum', sum('total').over(window)) \
                    .withColumn('Total_Count', count('total').over(window)) \
                    .orderBy("Total_Count") \
                    .show(5)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time	Total_Sum	Total_Count
91945	3	Jendouba_Ain_Drahem	Chocolate Eclair	6.0	7.5	45.0	livraison	online	2023-07-25 07:05:59	29970.0	399
51955	3	Jendouba_Ain_Drahem	Chocolate Eclair	15.0	7.5	112.5	direct	cash	2023-08-12 22:55:16	29970.0	399
57811	1	Jendouba_Center	Chocolate Eclair	9.0	7.5	67.5	direct	cash	2023-07-05 18:16:34	29970.0	399
43892	4	Jendouba_Tabarka	Chocolate Eclair	1.0	7.5	7.5	livraison	online	2024-01-05 06:53:42	29970.0	399
94275	2	Jendouba_Bousalem	Chocolate Eclair	16.0	7.5	120.0	livraison	online	2023-11-18 22:06:52	29970.0	399

only showing top 5 rows

Window + Ranking Functions

- row\_number()** - unique number for each row within partition, with different numbers for tied values
- rank()** - ranking within partition, with gaps and same ranking for tied values
- dense\_rank()** - ranking within partition, with no gaps and same ranking for tied values

city	price	row_number	rank	dense_rank
		over(order by price)		
Paris	7	1	1	1
Rome	7	2	1	1
London	8.5	3	3	2
Berlin	8.5	4	3	2
Moscow	9	5	5	3
Madrid	10	6	6	4
Oslo	10	7	6	4

Example 1

```
In [ ]: from pyspark.sql.window import Window
        from pyspark.sql.functions import row_number, desc
```

```
In [ ]: window = Window.partitionBy('article').orderBy('total')
        ranked_df = new_df.withColumn('rank', row_number().over(window))
        ranked_df.show(5)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time	rank
42619	2	Jendouba_Bousalem	Apple Turnover	1.0	5.0	5.0	direct	cash	2023-03-03 23:44:35	1
3649	3	Jendouba_Ain_Drahem	Apple Turnover	1.0	5.0	5.0	livraison	online	2023-04-30 08:45:21	2
11902	4	Jendouba_Tabarka	Apple Turnover	1.0	5.0	5.0	direct	cash	2023-04-30 10:18:21	3
44116	1	Jendouba_Center	Apple Turnover	1.0	5.0	5.0	livraison	online	2023-12-31 08:06:08	4
64700	3	Jendouba_Ain_Drahem	Apple Turnover	1.0	5.0	5.0	direct	card	2023-08-25 01:11:10	5

only showing top 5 rows

Example 2

```
In [ ]: window = Window.partitionBy('Article').orderBy(desc('Total'))
        ranked_df = new_df.withColumn('rank', row_number().over(window))
        ranked_df.show(5)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time	rank
2365	2	Jendouba_Bousalem	Apple Turnover	20.0	5.0	100.0	direct	card	2023-02-19 12:06:11	1
75660	4	Jendouba_Tabarka	Apple Turnover	20.0	5.0	100.0	livraison	online	2023-05-18 01:27:22	2
34432	1	Jendouba_Center	Apple Turnover	20.0	5.0	100.0	livraison	online	2023-09-05 01:10:37	3
76567	2	Jendouba_Bousalem	Apple Turnover	20.0	5.0	100.0	direct	card	2023-05-16 23:32:05	4
26798	2	Jendouba_Bousalem	Apple Turnover	20.0	5.0	100.0	direct	cash	2023-06-28 15:23:30	5

only showing top 5 rows

Window + Distribution Functions

- **cume\_dist()** the cumulative distribution of a value within a group of values, i.e., the number of rows with values less than or equal to the current row's value divided by the total number of rows; a value in (0, 1] interval
- **percent\_rank()** the percentile ranking number of a row—a value in [0, 1] interval: (rank-1) / (total number of rows - 1)

cume\_dist() OVER(ORDER BY sold)

city	sold	cume_dist
Paris	100	0.2
Berlin	150	0.4
Rome	200	0.8
Moscow	200	0.8
London	300	1

← 80% of values are less than or equal to this one

percent\_rank() OVER(ORDER BY sold)

city	sold	percent_rank
Paris	100	0
Berlin	150	0.25
Rome	200	0.5
Moscow	200	0.5
London	300	1

← without this row 50% of values are less than this row's value

```
In [ ]: from pyspark.sql.window import Window
from pyspark.sql.functions import col, cume_dist

windowSpec = Window.orderBy(col("total"))

cume_df = new_df.withColumn("cumulative_distribution", cume_dist().over(windowSpec))
cume_df.show(5)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time	cumulative_distribution
64726	3	Jendouba_Ain_Drahem	Croissant	1.0	1.5	1.5	livraison	online	2023-01-26 01:39:01	0.002103934357248054
82657	4	Jendouba_Tabarka	Croissant	1.0	1.5	1.5	direct	card	2023-02-09 21:36:48	0.002103934357248054
28496	4	Jendouba_Tabarka	Croissant	1.0	1.5	1.5	livraison	online	2023-08-30 16:56:55	0.002103934357248054
13206	2	Jendouba_Bousalem	Croissant	1.0	1.5	1.5	livraison	online	2024-01-16 03:17:05	0.002103934357248054
32696	3	Jendouba_Ain_Drahem	Croissant	1.0	1.5	1.5	livraison	online	2024-03-12 19:29:30	0.002103934357248054

only showing top 5 rows

Window + Analytic Functions

- **lead(expr, offset, default)** the percentile ranking number of a row—a value in [0, 1] interval: (rank-1) / (total number of rows - 1)
- **lag(expr, offset, default)** the cumulative distribution of a value within a group of values, i.e., the number of rows with values less than or equal to the current row's value divided by the total number of rows; a value in (0, 1] interval

lag(sold) OVER(ORDER BY month)

month	sold	
1	500	NULL
2	300	500
3	400	300
4	100	400
5	500	100

lead(sold) OVER(ORDER BY month)

month	sold	
1	500	300
2	300	400
3	400	100
4	100	500
5	500	NULL

lag(sold, 2, 0) OVER(ORDER BY month)

month	sold	
1	500	0
2	300	0
3	400	500
4	100	300
5	500	400

offset=2

lead(sold, 2, 0) OVER(ORDER BY month)

month	sold	
1	500	400
2	300	100
3	400	500
4	100	0
5	500	0

offset=2

- **first\_value(expr)** the value for the first row within the window frame
- **last\_value(expr)** the value for the last row within the window frame

first\_value(sold) OVER (PARTITION BY city ORDER BY month)

city	month	sold	first_value
Paris	1	500	500
Paris	2	300	500
Paris	3	400	500
Rome	2	200	200
Rome	3	300	200
Rome	4	500	200

last\_value(sold) OVER (PARTITION BY city ORDER BY month RANGE BETWEEN UNBOUNDED PRECEDING AND UNBOUNDED FOLLOWING)

city	month	sold	last_value
Paris	1	500	400
Paris	2	300	400
Paris	3	400	400
Rome	2	200	500
Rome	3	300	500
Rome	4	500	500

- **ntile(n)** divide rows within a partition as equally as possible into n groups, and assign each row its group number.

ntile(3)		
city	sold	
Rome	100	1
Paris	100	
London	200	
Moscow	200	2
Berlin	200	
Madrid	300	
Oslo	300	3
Dublin	300	

```
In [ ]: from pyspark.sql.window import Window
from pyspark.sql.functions import col, lag

windowSpec = Window.orderBy(col("sale_time"))

lag_df = new_df.withColumn("Previous Total", lag("Total", 1).over(windowSpec))
lag_df.show(5)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time	Previous Total
46899	1	Jendouba_Center	Palmier	9.0	3.0	27.0	direct	card	2023-01-01 00:48:33	null
46899	1	Jendouba_Center	Cream Puff	2.0	9.0	18.0	livraison	online	2023-01-01 00:48:33	27.0
46899	1	Jendouba_Center	Cream Puff	3.0	9.0	27.0	direct	cash	2023-01-01 00:48:33	18.0
70879	1	Jendouba_Center	Danish Pastry	12.0	6.5	78.0	direct	card	2023-01-01 06:22:22	27.0
70879	1	Jendouba_Center	Baguette	6.0	2.0	12.0	direct	cash	2023-01-01 06:22:22	78.0

only showing top 5 rows

## User-Defined Functions (UDF)

```
In [ ]: from pyspark.sql.functions import udf
from pyspark.sql.types import StringType, DoubleType
```

### Example 1

```
In [ ]: def kg_to_pounds(quantity):
    return quantity * 2.20462

kg_to_pounds_udf = udf(kg_to_pounds, DoubleType())

df_with_pounds = new_df.withColumn('Quantity_in_Pounds', kg_to_pounds_udf('Quantity'))
df_with_pounds.show(5)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time	Quantity_in_Pounds
15526	3	Jendouba_Ain_Drahem	Blueberry Muffin	12.0	3.8	45.6	livraison	online	2024-02-19 13:31:42	26.455439999999996
15526	3	Jendouba_Ain_Drahem	Bear Claw	20.0	6.8	136.0	livraison	online	2024-02-19 13:31:42	44.0924
15526	3	Jendouba_Ain_Drahem	Baguette	14.0	2.0	28.0	direct	card	2024-02-19 13:31:42	30.864679999999996
15526	3	Jendouba_Ain_Drahem	Lemon Bar	15.0	6.0	90.0	livraison	online	2024-02-19 13:31:42	33.0693
75376	3	Jendouba_Ain_Drahem	Muffin	14.0	3.5	49.0	direct	card	2023-03-08 18:29:41	30.864679999999996

only showing top 5 rows

### Example 2

```
In [ ]: def categorize_quantity(quantity):
    if quantity < 5:
        return 'Low'
    elif quantity >= 5 and quantity < 10:
        return 'Medium'
    else:
        return 'High'

categorize_udf = udf(categorize_quantity, StringType())

categorized_df = new_df.withColumn('Quantity_Category', categorize_udf(df['Quantity']))
categorized_df.show(5)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time	Quantity_Category
15526	3	Jendouba_Ain_Drahem	Blueberry Muffin	12.0	3.8	45.6	livraison	online	2024-02-19 13:31:42	High
15526	3	Jendouba_Ain_Drahem	Bear Claw	20.0	6.8	136.0	livraison	online	2024-02-19 13:31:42	High
15526	3	Jendouba_Ain_Drahem	Baguette	14.0	2.0	28.0	direct	card	2024-02-19 13:31:42	High
15526	3	Jendouba_Ain_Drahem	Lemon Bar	15.0	6.0	90.0	livraison	online	2024-02-19 13:31:42	High
75376	3	Jendouba_Ain_Drahem	Muffin	14.0	3.5	49.0	direct	card	2023-03-08 18:29:41	High

only showing top 5 rows

# Spark SQL

```
In [ ]: new_df.createOrReplaceTempView("sales")
```

## Example 1: Show Data

```
In [ ]: query1 = """
SELECT *
FROM sales
"""

result1 = spark.sql(query1)
result1.show(5)
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time
15526	3	Jendouba_Ain_Drahem	Blueberry Muffin	12.0	3.8	45.6	livraison	online	2024-02-19 13:31:42
15526	3	Jendouba_Ain_Drahem	Bear Claw	20.0	6.8	136.0	livraison	online	2024-02-19 13:31:42
15526	3	Jendouba_Ain_Drahem	Baguette	14.0	2.0	28.0	direct	card	2024-02-19 13:31:42
15526	3	Jendouba_Ain_Drahem	Lemon Bar	15.0	6.0	90.0	livraison	online	2024-02-19 13:31:42
75376	3	Jendouba_Ain_Drahem	Muffin	14.0	3.5	49.0	direct	card	2023-03-08 18:29:41

only showing top 5 rows

## Example 2: SUM Function

```
In [ ]: query2 = """
SELECT SUM(total) AS total_sales
FROM sales
"""

result2 = spark.sql(query2)
result2.show()
```

total_sales
642509.0000000019

## Example 3: Where Condition

```
In [ ]: query3 = """
SELECT article, quantity
FROM sales
WHERE Quantity > 5
"""

result3 = spark.sql(query3)
result3.show(5)
```

article	quantity
Blueberry Muffin	12.0
Bear Claw	20.0
Baguette	14.0
Lemon Bar	15.0
Muffin	14.0

only showing top 5 rows

## Example 4: Group By

```
In [ ]: query4 = """
SELECT article, AVG(total) AS average_total
FROM sales
GROUP BY article
"""

result4 = spark.sql(query4)
result4.show(5)
```

article	average_total
Cream Puff	94.9691943127962
Muffin	37.177033492822964
Pecan Pie	115.92098765432098
Napoleon	82.37739557739553
Scone	29.338596491228053

only showing top 5 rows



# Python Data Visualization



```
In [ ]: import matplotlib.pyplot as plt
import pandas as pd
```

```
In [ ]: new_df.show(3)
new_df.count()
```

client_id	pos_id	pos_name	article	quantity	price	total	sale_type	payment_mode	sale_time
15526	3	Jendouba_Ain_Drahem	Blueberry Muffin	12.0	3.8	45.6	livraison	online	2024-02-19 13:31:42
15526	3	Jendouba_Ain_Drahem	Bear Claw	20.0	6.8	136.0	livraison	online	2024-02-19 13:31:42
15526	3	Jendouba_Ain_Drahem	Baguette	14.0	2.0	28.0	direct	card	2024-02-19 13:31:42

only showing top 3 rows

Out[ ]: 9506

## Example 1: Donut (Pie) Chart

```
In [ ]: from pyspark.sql.functions import sum

article_totals = new_df.groupBy('article').agg(sum('total').alias('Total')).orderBy('Total', ascending=False)


article_totals_pd = article_totals.toPandas()
```

```
In [ ]: article_totals_pd = article_totals_pd[:5]
```

```
In [ ]: plt.figure(figsize=(8, 8))
plt.pie(article_totals_pd['Total'], labels=article_totals_pd['article'], autopct='%1.1f%%', startangle=90)

centre_circle = plt.Circle((0,0),0.70,fc='white')
fig = plt.gcf()
fig.gca().add_artist(centre_circle)

plt.axis('equal')
plt.title('Total Sales by Article')
plt.tight_layout()
plt.show()
```



## Example 2: Bar Chart

```
In [ ]: from pyspark.sql.functions import sum

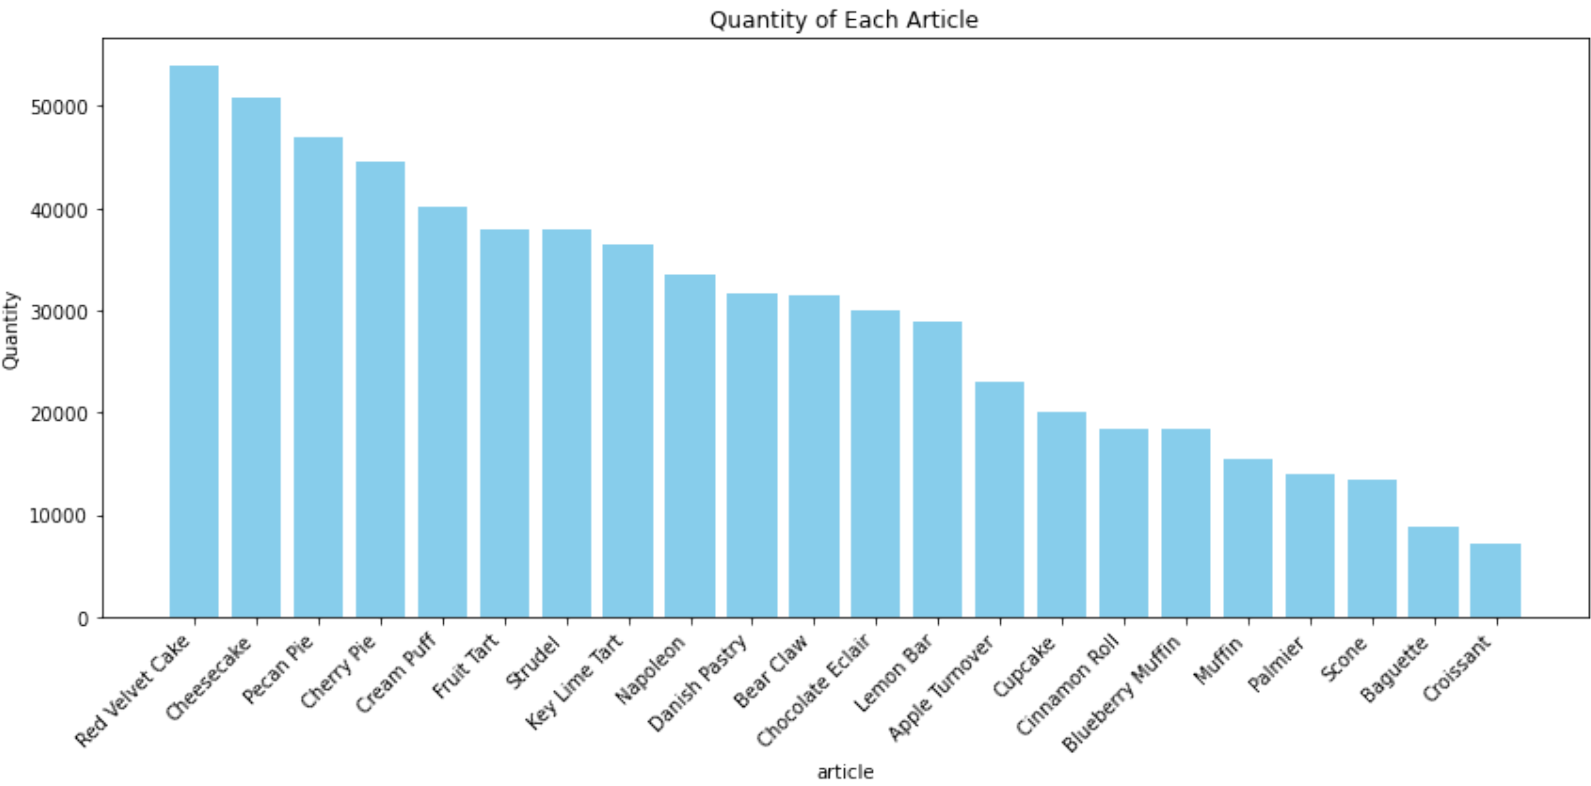
article_quantity = new_df.groupBy('article').agg(sum('quantity').alias('Total')).orderBy('Total', ascending=False)
article_quantity_pd = article_totals.toPandas()

In [ ]: article_quantity_pd.head(5)
```

Out[ ]:

	article	Total
0	Red Velvet Cake	53940.0
1	Cheesecake	50778.0
2	Pecan Pie	46948.0
3	Cherry Pie	44517.0
4	Cream Puff	40077.0

```
In [ ]: plt.figure(figsize=(12, 6))
plt.bar(article_quantity_pd['article'], article_quantity_pd['Total'], color='skyblue')
plt.title('Quantity of Each Article')
plt.xlabel('article')
plt.ylabel('Quantity')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()
```



## Example 3: Line Chart

```
In [ ]: df_Croissant = new_df.where(df['article'] == "Croissant") \
        .select("article", "quantity", "sale_time") \
        .orderBy("sale_time", ascending=True) \
        .dropDuplicates(["sale_time"])

df_Croissant_pd = df_Croissant.toPandas()
df_Croissant_pd.set_index('sale_time', inplace=True)

In [ ]: df_Croissant_pd.head(5)
```

Out[ ]:

	article	quantity
sale_time		
2023-01-01 06:22:22	Croissant	17.0
2023-01-02 01:59:10	Croissant	15.0
2023-01-04 08:39:15	Croissant	9.0
2023-01-04 13:20:38	Croissant	11.0
2023-01-04 20:48:10	Croissant	9.0

```
In [ ]: plt.figure(figsize=(20, 8))
plt.plot(df_Croissant_pd.index, df_Croissant_pd['quantity'])
plt.title('Quantity vs. Sale Time')
plt.xlabel('sale_time')
plt.ylabel('quantity')
plt.xticks(rotation=45)
plt.grid(True)
plt.tight_layout()
plt.show()
```

